

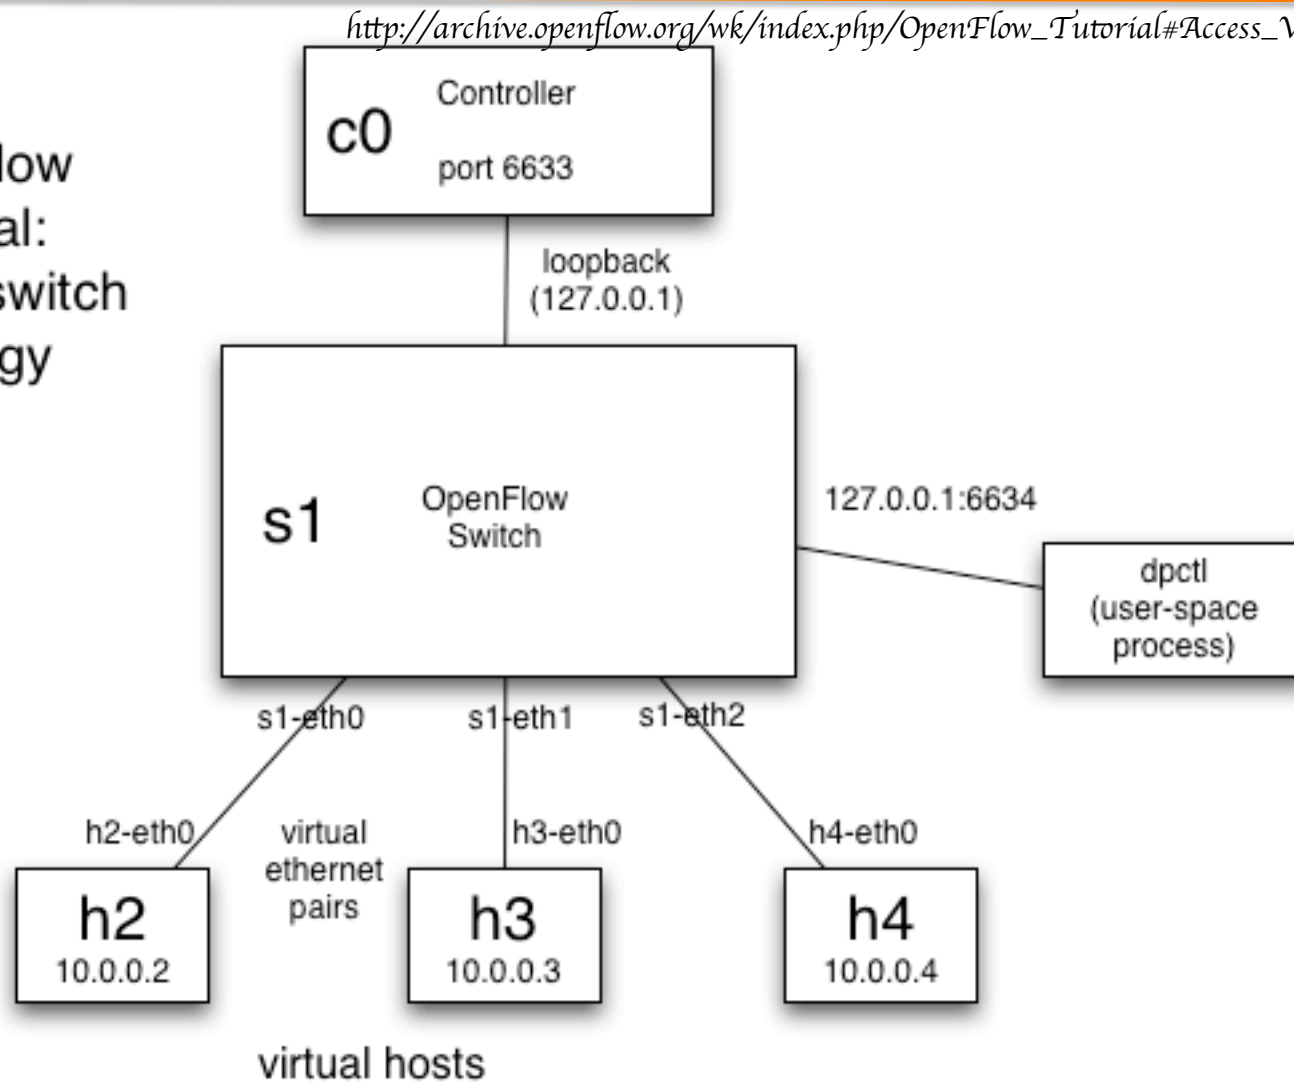
MININET BASED STUDIES

Mininet Demo

- “Mininet creates a **realistic virtual network**, running **real kernel, switch and application code**, on a single machine (VM, cloud or native”
– mininet.org
 - Interact with network using CLI or API
 - Useful for experiments with OpenFlow/SDN
 - BSD Open Source license
- Mininet VM is a good way to get started
 - Can also build on Linux
- Can interact with SDN controllers (Floodlight, POX, etc.)


Mininet Topology

OpenFlow
Tutorial:
3hosts-1switch
topology




Different Topologies


Mininet Topology



Single Topology



Tree Topology



Linear Topology

POX Controller Programming

- Python interface
 - Write scripts in Python to program controller
- When a connection to a switch starts, a ConnectionUp event is fired
- **connection.send()**
 - To send OF messages from controller to switch
- **ofp_action_output class**
 - `out_action = of.ofp_action_output(port = of.OFPP_FLOOD)`
 - `out_action = of.ofp_action_output(port = 2)`

POX Documentation

- <https://openflow.stanford.edu/display/ONL/POX+Wiki>

Launch

```
def launch (): # Definition Start
```

```
    def start_switch (event):  
        L2Hub(event.connection)
```

```
core.openflow.addListenerByName("ConnectionUp", start_switch) # Definition End
```

Parsing packets

```
packet = event.parsed
src_mac = packet.src
dst_mac = packet.dst
if packet.type == ethernet.IP_TYPE:
    ipv4_packet = event.parsed.find("ipv4")
    # Do more processing of the IPv4 packet
    src_ip = ipv4_packet.srcip
    dst_ip = ipv4_packet.dstip
```


Packets parsed by pox/lib

- arp, dhcp, dns
- eapol, eap
- ethernet, icmp
- igmp, ipv4
- llc, lldp, mpls
- rip
- tcp, udp, vlan

Sending message to switch

```
msg = of.ofp_packet_out()          # Create packet out message
msg.buffer_id = event.ofp.buffer_id # Use the incoming packet as the
data for the packet out
```

```
msg.in_port = packet_in.in_port    # Set the in_port so that the switch
knows
```

```
# Set match fields
msg.match = of.ofp_match.from_packet(packet)
```

```
# Add an action to send to the specified port
```

```
action = of.ofp_action_output(port = of.OFPP_FLOOD)
msg.actions.append(action)
```

```
# Send message to switch
self.connection.send(msg)
```

Events gen. from switch message

- FlowRemoved, FeaturesReceived
ConnectionUp, FeaturesReceived
RawStatsReply, PortStatus
PacketIn, BarrierIn, SwitchDescReceived,
FlowStatsReceived, QueueStatsReceived,
AggregateFlowStatsReceived,
TableStatsReceived, PortStatsReceived

POX Programming, contd.

➤ **ofp_match class**

- Objects specify packet header fields and input port to match on
 - dl_src - The data link layer (MAC) source address
 - dl_dst - The data link layer (MAC) destination address
 - in_port - The packet input switch port
- All fields are optional
- Example:
 - match = of.ofp_match()
 - match.in_port = 3

POX Prog., contd.

➤ **ofp_packet_out OpenFlow message**

- instructs a switch to send a packet
- Packet constructed by switch (OR) packet earlier switch by switch to Controller for processing
- Example

```
msg = of.ofp_packet_out()  
newaction = of.ofp_action_output(port = out_port)  
msg.actions.append(newaction)
```

```
# Send message to switch  
self.connection.send(msg)
```

POX Prog., contd.

➤ **ofp_flow_mod OpenFlow message**

- instructs a switch to install a flow table entry
- Entries match packet fields and execute actions
- Components of entry
 - idle_timeout - Number of idle seconds before the flow entry is removed. Defaults to no idle timeout.
 - hard_timeout - Number of seconds before the flow entry is removed. Defaults to no timeout.
 - actions - A list of actions to perform on matching packets (e.g., ofp_action_output)
 - priority - When using non-exact (wildcarded) matches, this specifies the priority for overlapping matches. Higher values are higher priority. Not important for exact or non-overlapping entries.
 - match - An ofp_match object. By default, this matches everything, so you should probably set some of its fields
- `fm = of.ofp_flow_mod()`
- `fm.match.in_port = 3`
- `fm.actions.append(of.ofp_action_output(port = 4))`

Mininet, Contd.

Switch Types

- User-space Switch (Slow)
- Kernel-space Open Virtual Switch (Fast)
- Hardware Switch supporting OpenFlow
- FPGA Switches


Controller

- Local Default Controller
- Local NOX/POX/Floodlight etc. Controller
- Remote Controller

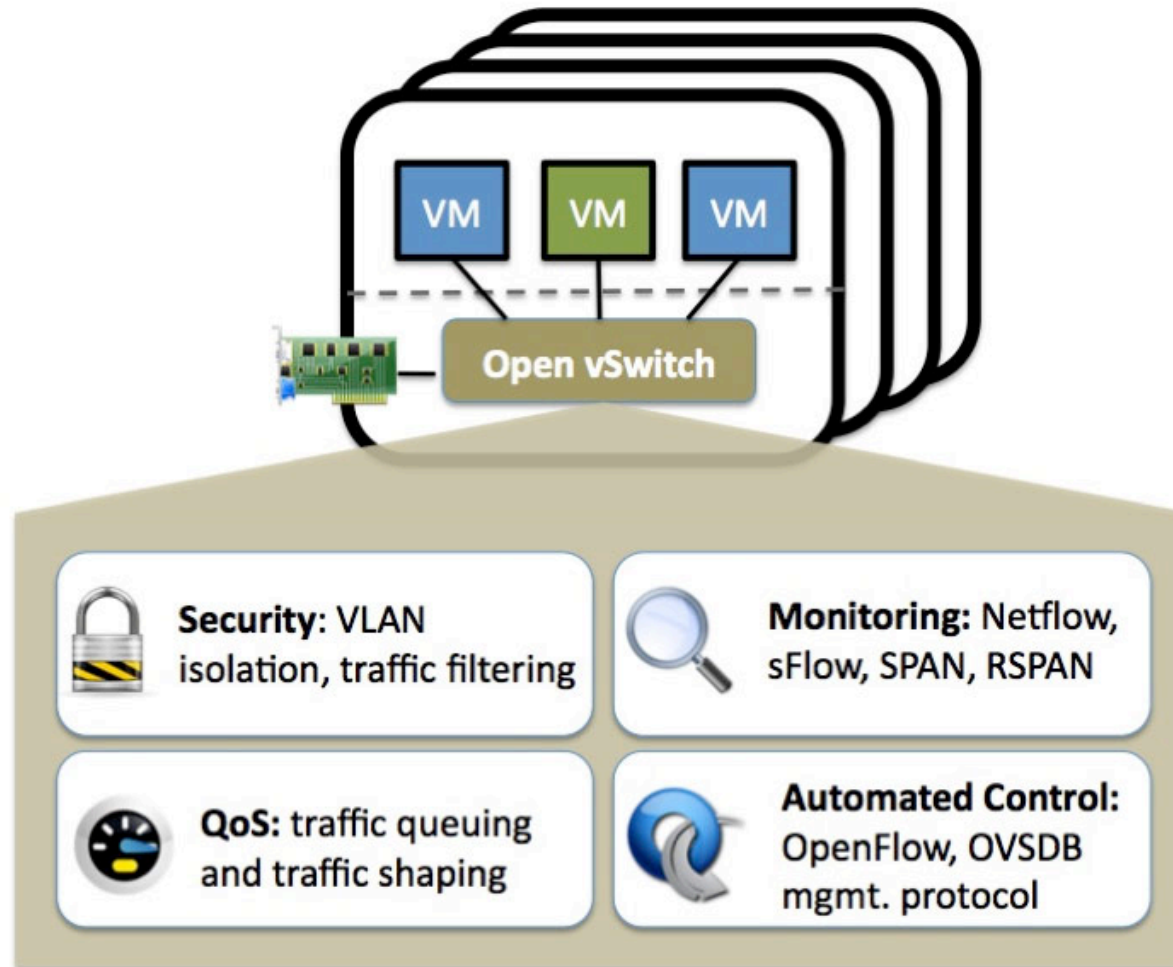
Open Virtual Switch (OVS)

- <http://openvswitch.org/>
- Apache 2.0 License
- “production quality, multilayer virtual switch”
- “Supports distribution across multiple physical servers similar to VMware's vNetwork distributed vswitch or Cisco's Nexus 1000V”
- Runs a soft switch inside the hypervisor, or as part of software stack of switching h/w
- Default switch in Xen Server, Xen Cloud
- Integrated into OpenStack, OpenNebula, oVirt
- Kernel datapath is distributed with [Linux](#), with packages for [Ubuntu](#), [Debian](#), and [Fedora](#)

OVS, conceptual rep.



OVS



Source: <https://www.sdxcentral.com/wp-content/uploads/2014/09/open-vswitch-sdn-virtualization.jpg>

NETWORK VIRTUALIZATION

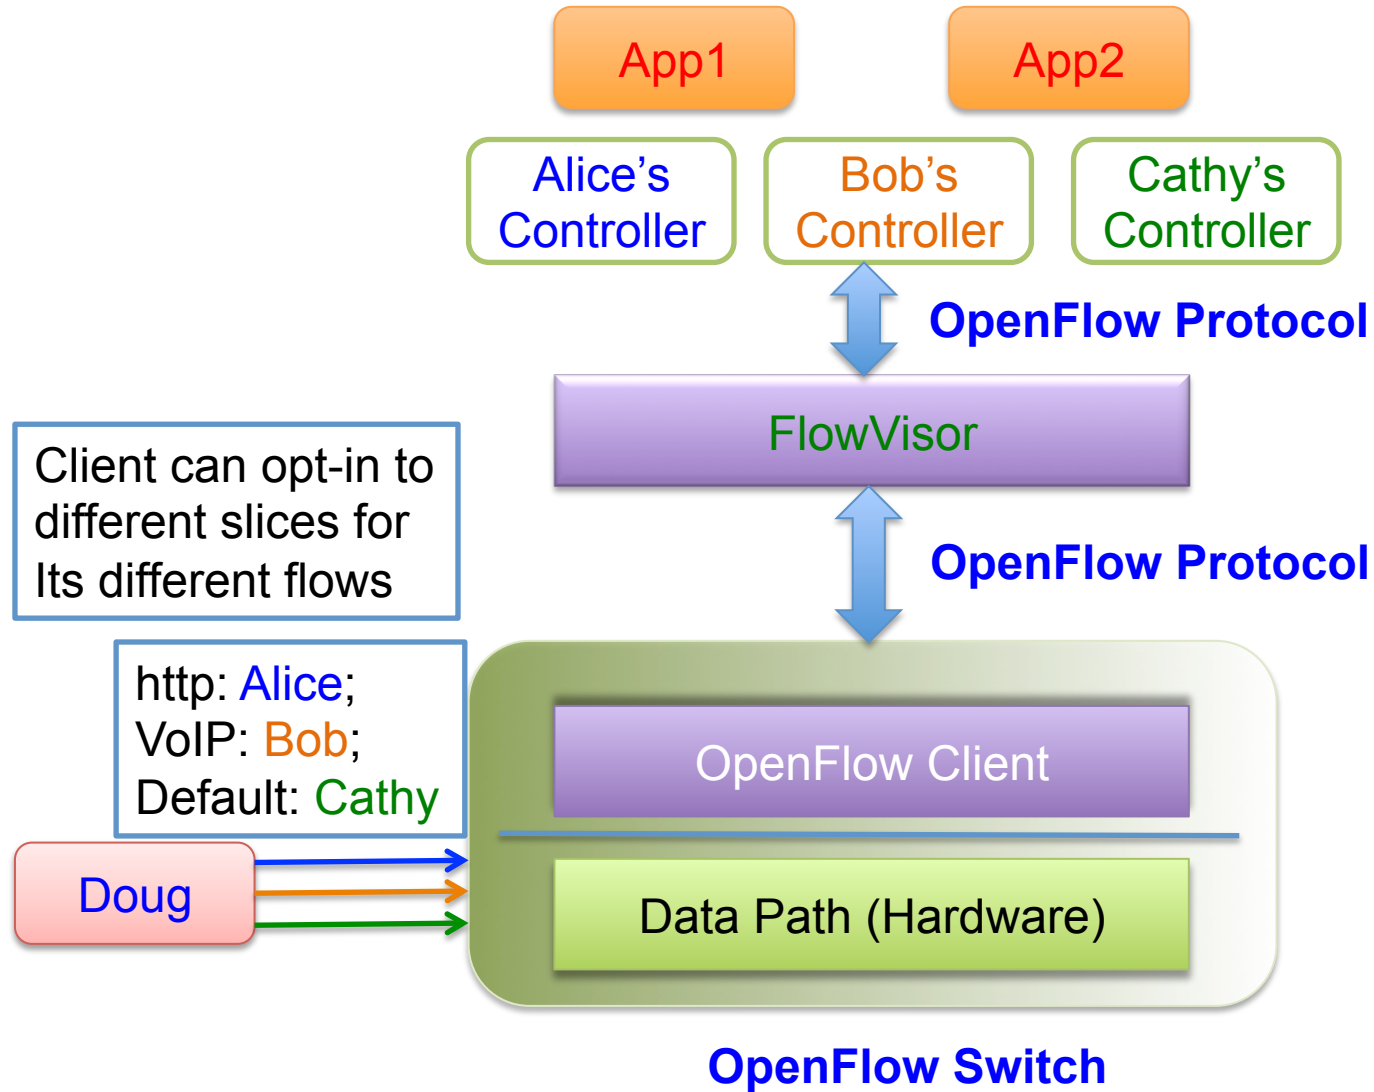
Network Virtualization

- Compute & Storage Virtualization in place
- Virtualization of Networks is up next
 - Going beyond VLANs
- Creation of “Virtual Network Elements” at different levels
 - Switches, Networks, etc.
- Virtualization Enabling Systems Available
 - FlowVisor, RouteFlow, VMware NSX, etc.

Stanford's FlowVisor

- Provides a software-based “slicing layer” between forwarding and control planes
- Policy language to map flows to slices
 - Each slice controls a set of flows, “flowspace”
 - Resources sliced in terms of: BW, Topology, Device CPU, Forward Table Entries, etc.
- Can operate on deployed networks
 - The production network can be a testbed
- Multiple Controllers exist above FlowVisor
 - Each controller uses its logic on its slice

Flow Visor, Contd.




FlowSpace


- Defined by collection of packet headers and assigned to “Slices”
 - Source/Destination MAC address
 - VLAN ID
 - Ethertype
 - IP protocol
 - Source/Destination IP address
 - Source/Destination port number
 - ToS/DSCP

Slices


Slice: OpenRoads




Slice: Aggregation




Slice: Production




Slice: PlugServ



Physical Network



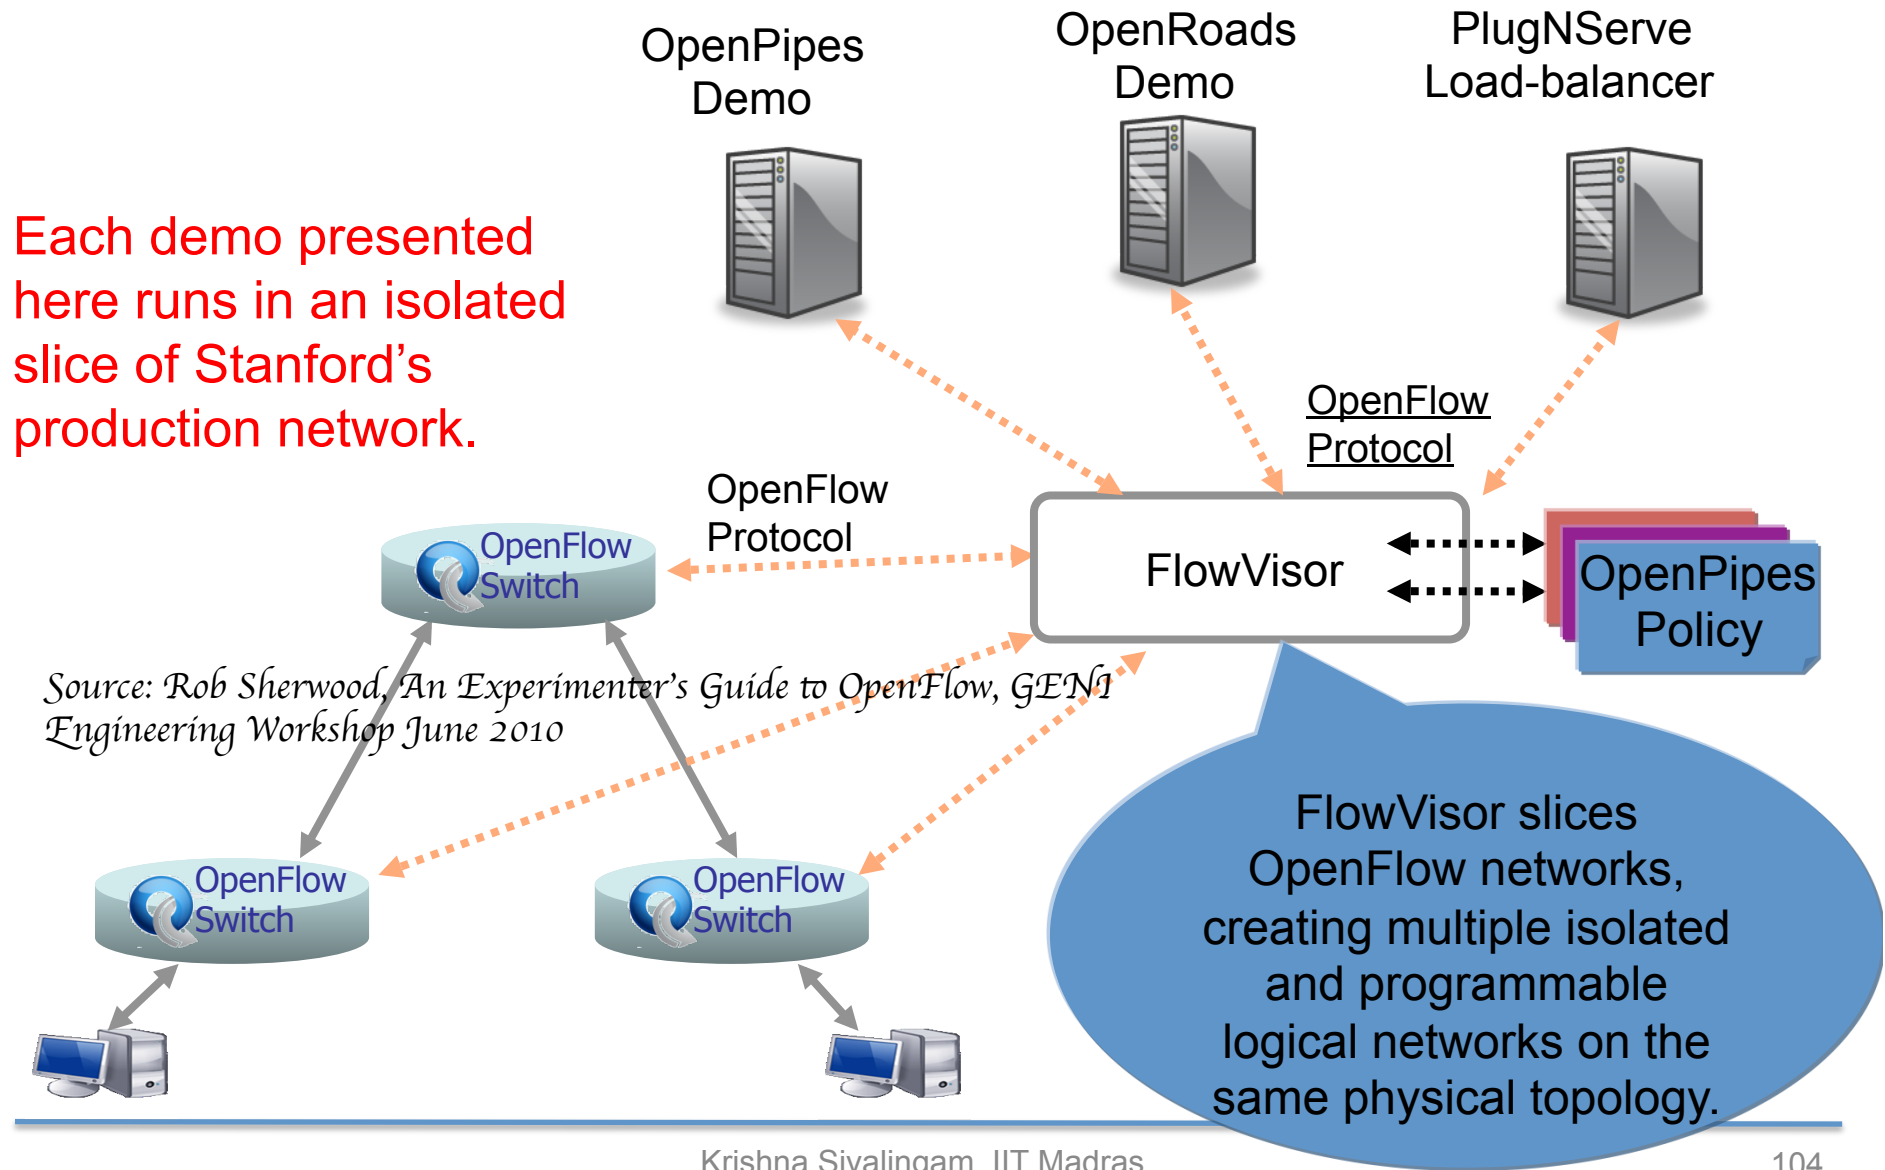
Slice: OpenPipes



Source: Rob Sherwood, *An Experimenter's Guide to OpenFlow, GENI Engineering Workshop June 2010*

FlowVisor Creates Virtual Networks

Each demo presented here runs in an isolated slice of Stanford's production network.



OpenVirteX (OVX): Network Virtualization

- From ON.Lab at Stanford
- Single Infrastructure provider
 - Multiple tenants, each specifies network and resource connectivity
- OVX supports multiple virtual networks
 - Helps specify tenant's topology and addressing
- OVX resides between physical hardware and virtual network controllers
 - Create isolated virtual networks
 - Use tenant's choice of Network OS
 - Use entire address space!


ON.LAB OVX Architecture



Network Operating
System Layer



Network Virtualization
Layer



Physical Infrastructure

VMWare NSX

- Provides network virtualization for software-defined data center
 - Based on Nicira Acquisition (2012)
- Treats physical network as a pool of transport capacity
- Virtual Networks can be created
 - Logical switches, routers, firewalls, load balancers, etc.
 - Programmatically created, provisioned and managed
- <https://www.vmware.com/products/nsx/features.html>